

An Introduction to High Performance Computing

Serial Assignment Feedback

8 November 2018

General Feedback

- A great start to this year's HPC course!
- Many of you clearly worked hard on your submissions
- Remember, the overall quality of your report is very important
- Many of you only tried the minimum expected optimisations
- Not everything here will apply to **you**

Report Feedback

Report Feedback

- State which compiler (and what version) you use
- Lack of explanation and analysis of results
 - “I did this and it was faster”... but **why**?!
 - Back up with evidence
- Don't report the same data twice (in a table **and** a graph)
- Analyse **all** the problems, not just one of them
 - They show different performance characteristics
 - The small one is the least interesting for HPC

Report Feedback

- Screenshots, terminal output, etc. are data. You must interpret and present this – no screenshots!
- Proofread – check spelling and grammar
- Page limit is **strict**, and must include everything, including references
 - Next time you **will lose marks**
 - No exceptions because you think you did more than was asked so you needed more space to talk about it
 - Many of you had lots of whitespace in...

Report Feedback

- Some of you didn't include a title, or even your name!
 - This was part of the spec!
- Use fixed-width (monospace) font for code
 - Otherwise text can get hard to read and it may not be obvious when you're referring back to the code
- When reporting a run time of 596 seconds, does it matter that it is actually 596.2186785 seconds?
 - Don't use more decimal places than necessary!
 - Remember there is noise, so there are error margins: 0.093 vs 0.098 s is not really a big improvement

Report Feedback

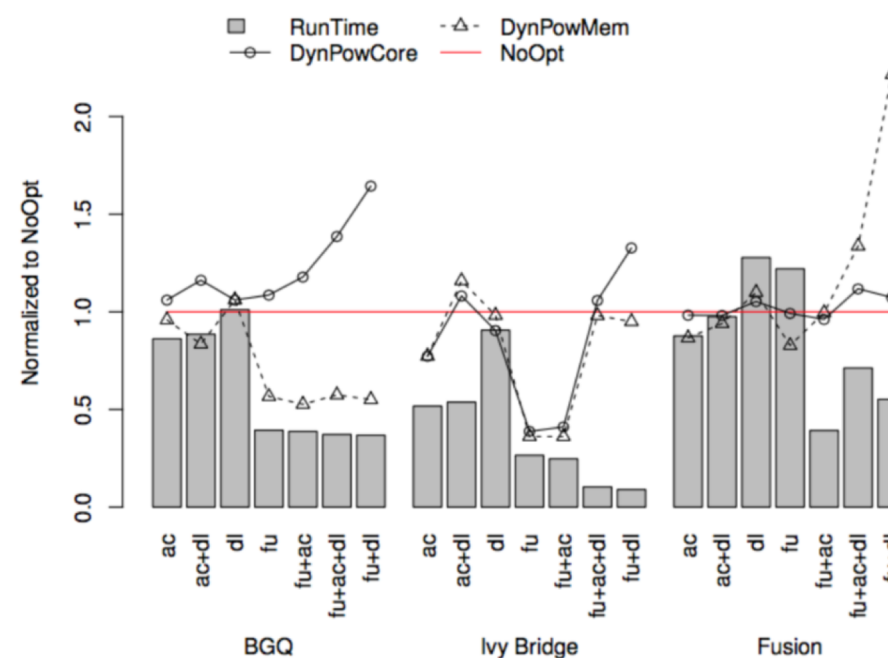
- Use subheadings to break up large sections of text
 - Don't use strange subtitles either, e.g. Body
 - Don't stuff a whole section into a single paragraph
- Put evidence, timings, speedups, etc. next to the relevant text; **not** in a summary table at the end
- But do include a summary table of final runtimes for all inputs
 - Make it clear what your final configuration was

Report Feedback

- Speedup as a % is confusing
 - 100% faster might be same speed, or half speed, or ...
- Better to use X times faster: $\text{old_time}/\text{new_time}$
 - E.g. 100 s to 50 s results in a 2X speedup
- Useful to measure your achieved memory bandwidth and arithmetic performance
 - Compare to the hardware's peak to see how much room for improvement there is
 - But keep in mind that *peak* may not be *achievable*

Report Feedback

- **Don't** do a timeline of optimisation graph (or table)
- Many optimisations depend on previous optimisations
- This graph is a good one, for only 3 optimisations, but would be intractable for your coursework



Common Mistakes

Common Mistakes

- CC compiler is the system default (GCC 4.4.7)
- Incorrect filenames and build errors
 - Check what you submit!
 - Be careful with whitespace issues
- Test **all** inputs
 - Must complete running and validate
 - Don't do your work based on a single one
- Next time you will lose **more marks** for build and run errors

Common Mistakes

- The provided code doesn't trash the cache
 - It just makes very poor use of it
- A lot of you mentioned that the compiler does plenty of optimisations...
 - ... but attempted to do them all by hand before using optimisation flags...
 - ... or used an old compiler version
 - You can tell if the compiler has done an optimisation you would have done by hand by looking at the assembly code
- The stencil code is **memory-bandwidth-bound**
 - Not memory-bound

Row/Column-Major

Lots of mistakes in this sentence!

- “C is a row-major language, so I changed from A being column-major to row-major”
- Arrays are on the heap and use pointer arithmetic to access data, so C stack arrays are irrelevant here
- Row-major Column-major



- What's important: ensure data order and access pattern are the **same**

Optimisation Feedback

Optimisation Feedback

- Start with changing the compiler and flags
 - It's no-effort optimisation!
 - Has a big effect which nullified other manual changes later
- Intel compiler defaults to -O2
 - Compare to GCC fairly by turning off optimisations: -O0
- If -O3 is not faster than -O2, then your code likely can't take advantage of some hardware features, e.g. vectorisation
 - Similarly, if you do an optimisation by hand and you don't observe the expected result, double-check that you haven't made an implementation mistake

Optimisation Feedback

- Vectorisation is **really** important for performance
 - Check the compiler vector reports `-qopt-report=5`
 - Many of you simply checked the loop vectorised at all, and didn't investigate further
 - Look at vector width, alignment, access pattern, types of load/stores, etc.
 - Did the most performance important loops vectorise, not just unimportant ones?

Optimisation Feedback

- Removal of branches in main loop:
 - Many of you said this was due to branch misprediction causing bubbles in the pipeline
 - Some of you said they cause branch predictions, which should be avoided
 - Nobody presented any evidence that this is an issue
 - The compiler might generate masked vector operations, or it might produce a peel loop
 - Your aim is to get the compiler to understand what you've done, so that it can enable **effective vectorisation**
 - Some of you noticed that the compiler sometimes understands you code better if you leave the branches in!

MPI Assignment

- Deadline is 5pm, Friday 14th December (week 11)
- Take our feedback into account (both individual and from today's lab)
- Start with your optimised serial code
 - Make sure it works on non-square grids, as the grid is unlikely to be decomposed to a square on every rank
- Do include any new serial optimisations you find, but don't repeat yourself from Assignment 1
- Remember, this is a report marked coursework too
 - Leave enough time to write it